Quiz 3

Question 1 (1 point)

S3 buckets have globally unique names yet have region affinity.

○ True

○ False

Question 2 (1 point)

Buckets and objects stored in buckets are addressable via an ARN

○ True

○ False

Question 3 (1 point)

Which are types of meta data you can attach to an S3 object

☐ Content-Encoding

☐ Content-Length

☐ Content-Owner

☐ Content-Type

Question 4 (1 point)

Eventual Consistency means

○ When the Elastic Load Balancer detects a faulty web server and removes it

○ When objects just written to an S3 bucket will be immediately available to all readers

○ When objects just written to an S3 bucket will be available to all readers soon

○ When instances in an EC2 scaling group eventually converge to the same runtime status

Question 5 (1 point)

Which factors contribute to the high level of durability of S3

☐ High level of replication between regions

☐ High level of <u>replication within a region</u>

☐ <u>Self-healing data</u>

☐ Being backed by SSDs

Question 6 (1 point)

If you are not sure of the usage pattern of your data and you want the best of both lower cost and high performance you should use

○ <u>Intelligent tiering</u>

○ EBS storage

○ Glacier storage

○ Standard storage

Question 7 (1 point)

The most straightforward encryption option for S3 is

○ RSA-56

○ SSE-KMS

○ SSE-C

○ SSE-S3

Question 8 (1 point)

Data before it is encrypted is called

○ Plaintext

○ Ciphertext

○ Data at rest

○ BinaryText

Question 9 (1 point)

Symmetric encryption requires 2 keys

○ True
○ False

Question 10 (1 point)

Which of the following is/are not a reason(s) to use RDS

☐ RDS enables reliable failover

☐ You have an app that needs a managed instance of MongoDB

☐ Letting AWS manage a production instance of is PostgreSQL cheaper than doing it yourself

☐ You like to shear yaks

Question 11 (1 point)

RDS enables automated snapshots as well as point-in-time restores

○ True
○ False

Question 12 (1 point)

In general, RDS enables infinite horizontal scaling

○ True

○ False

Question 13 (1 point)

A read replica is

○ A crude way to scale vertically

○ A copy of a database kept in sync with the primary database used to lessen the load of the primary database

○ Used to improve the write performance of a database

○ A snapshot of a database at a point in time that you can generate reports from

Question 14 (1 point)

Multi-AZ database deployments enable database software updates and other maintenance tasks to happen with minimal downtime

○ True

○ False

Question 15 (1 point)

You should be careful what you encrypt at rest because it can impact performance

○ True

○ False

Question 16 (1 point)

The only RDS hosted database that is serverless and scales horizontally well is

○ MongoDB

○ Oracle

○ Aurora

○ PostgreSQL

Question 17 (1 point)

NoSQL means Not Only SQL

○ True
○ False

Question 18 (1 point)

Relational databases are overkill

○ True
○ False

Question 19 (1 point)

A key weakness of NoSQL databases is that they do not support transactions

○ True
○ False

Question 20 (1 point)

A key strength of DynamoDB is its durability and its performance

○ True
○ False

Question 21 (1 point)

Which of the following are features of DynamoDB

☐ Transactions

☐ Automatic multi-region replication

☐ Infinite scalability

☐ Unlimited Secondary Indexes

Question 22 (1 point)

DynamoDB **query** operation is faster than **scan**

○ True
○ False

Question 23 (1 point)

TTL is used to

○ Automatically expire DynamoDB records after a period of time

○ Automatically take snapshots of dynamoDB tables after a period of time

○ Automatically log DynamoDB usage data to CloudTrail

○ Benchmark DynamoDB's query operations

Question 24 (1 point)

DynamoDB may be fast, but caching is still a useful performance improving strategy for DynamoDB

○ True
○ False

Submit Quiz   *0 of 24 questions saved*